

LAUNCH PLAN • CURRENT STATUS • FUTURE PLANS

# Let agents spend. *Keep the control.*

Fiducia controls how AI agents spend money. We start with one job: **research agents that buy datasets, APIs, reports and specialist work for an agency's clients.** Every payment runs inside an on-chain mandate on Solana, every stage waits for a human, and the seller is paid only for what was delivered.



## 33%

of enterprise software  
apps expected to include  
agentic AI by 2028, up  
from under 1% in 2024

Gartner, Oct 2024

## 40%+

of agentic AI projects  
predicted to be cancelled  
by end 2027, citing cost,  
value and risk controls

Gartner, Jun 2025

## \$150B+

global insights industry:  
the spend pool research  
agents draw from

ESOMAR Global Market Research  
2024, as reported

## ~\$0.0009




median Solana transaction  
fee, so a 5 cent data call  
can settle on chain

Blockworks Research, Q4 2025

External figures are market context only. They are not Fiducia traction, revenue or obtainable market. Fiducia today: live Solana devnet product, open-source repo, no paying customers yet.

# Agents are getting budgets. *Controls are not keeping up.*

The demand signal is not "agents can pay". It is that companies are deploying agents faster than they can govern them, and the projects that fail cite exactly the gap Fiducia closes: cost and risk control.



FIVE DATED SIGNALS

Apr 2025

Visa Intelligent Commerce and Mastercard Agent Pay announced: card networks are preparing for agent-initiated payments.

May 2025

Coinbase launches x402, HTTP-native stablecoin payments for APIs and agents.

Sep 2025

Google's AP2 and OpenAI/Stripe's Agentic Commerce Protocol: both rely on user "mandates" or delegated authority.

Oct 2025

McKinsey: agentic commerce could mediate \$3T to \$5T of global consumer commerce by 2030.

Q4 2025

Solana: ~\$14.9B stablecoin supply and sub-cent median fees (Blockworks).

The rails are arriving. None of them decides *how much* one agent may spend, *with whom*, at *which stage*, and whether the work was delivered. That layer is Fiducia.

# Built first for *teams that deploy research agents*.

IDEAL FIRST CUSTOMER

Small AI product teams and automation agencies (2 to 30 people) that run research agents for clients.

- ◆ **Buyer:** technical founder or lead developer.
- ◆ **Pain:** someone answers to a client for every dollar an agent spends.
- ◆ **Need:** caps, approved sellers, human checkpoints, per-client audit trail.

A CONCRETE ACCOUNT

An agency runs a market-research agent for **10 clients**. Each client's agent buys pricing feeds, company data, reports and expert reviews every week.

Today: shared API keys, a company card, and a spreadsheet at month end. One bad prompt can buy the wrong dataset for the wrong client.

**With Fiducia:** one mission and one mandate per client, payees fixed in advance, every stage approved, every receipt on chain.

HOW BIG IS THE FIELD

**976,300**  
U.S. professional, scientific and technical services establishments (NAICS 54)  
U.S. Census, 2022 Economic Census First Look

**\$4.0B**  
gross services volume through Upwork in 2025  
Upwork 10-K, FY2025

WHY RESEARCH PROCUREMENT FIRST: OUR SCORING OF CANDIDATE VERTICALS

| VERTICAL             | PURCHASES PER TASK   | TICKET SIZE     | SELLER IS KNOWN | DELIVERY IS HASHABLE | ACCOUNTABLE TO A CLIENT | REGULATORY DRAG       | VERDICT    |
|----------------------|----------------------|-----------------|-----------------|----------------------|-------------------------|-----------------------|------------|
| Research procurement | many                 | \$0.01 to \$50  | yes             | yes: file, report    | yes                     | low                   | START HERE |
| Coding agents        | some (APIs, compute) | \$0.01 to \$20  | yes             | yes: diff, build     | sometimes               | low                   | NEXT       |
| Travel and booking   | few                  | \$50 to \$2,000 | mixed           | partly               | yes                     | medium                | LATER      |
| Consumer commerce    | few                  | \$5 to \$500    | no              | no: physical         | no                      | high (cards, refunds) | LATER      |

Our qualitative assessment. Research agents make frequent small purchases from sellers that can be named in advance, and the output is a file whose hash can be checked, which is exactly what an on-chain mandate and escrow enforce best.

FOUR FIRST USE CASES

| JOB                       | AGENT BUYS                | MANDATE EXAMPLE                        |
|---------------------------|---------------------------|----------------------------------------|
| Competitor pricing report | pricing feed, 2 reports   | 8 USDC total, 3 per payment, 3 sellers |
| Due-diligence brief       | company data, expert call | 40 USDC, stage 2 needs approval        |
| Literature scan           | paper access, summaries   | 5 USDC, 0.25 per call (x402)           |
| Market map for a client   | datasets, analyst review  | 25 USDC, analyst paid on delivery      |

Illustrative mandates for design-partner pilots, not live customers.


WHAT THE WINNERS TAUGHT US

- ◆ **Superteam × WHU 2026** (we placed 3rd): 1st and 2nd named one local user and one measured outcome. We led with the mechanism.
- ◆ **ETHOnline 2026:** about 60 of 812 projects built agent spending caps; caps alone did not place. Finalists in our space (OpenBook, OnchainRouter, Cordon) sold a buyer story: data with refunds, pay per answer, team budgets.
- ◆ **Gap none of them covered:** per-client accountability for agencies, enforced on Solana, with escrow on delivery. That is our wedge.

Our own analysis of public showcases and finale, Oct 2026.

## Two sides, *one set of rules*.

Buyer agents spend inside programmable rules. On the other side, anyone with something an agent needs can sell it: data owners, API makers, researchers, developers selling agents, and architects selling designs or specialist reviews. The same escrow protects both.



### SPECIMEN 01 • MANDATE

Total cap, per-payment cap, approved sellers, stages and expiry for each agent. Enforced by the Solana program, so a buggy or hijacked agent cannot exceed it.

### SPECIMEN 02 • STAGE APPROVAL

The human approves the exact plan hash before any money moves in that stage. One click to revoke any agent.


### SPECIMEN 03 • ESCROW

The seller or agent team is paid only for the delivered product hash, with a review window, challenge and refund path.

### SUPPLY STRATEGY: SEED THE SHELF BEFORE THE BUYERS ARRIVE

- ◆ **Hand-pick 30 research listings** for the private beta: public datasets repackaged with clear terms, pricing feeds, analyst templates. Each one graded and priced by the seller pipeline already in the repo.
- ◆ **Zero listing fee and zero seller commission** at launch; the fee is charged on the buyer side when a payment settles.
- ◆ **Demand board** (already an MCP tool): agents post what they could not find; we recruit sellers for the top requests.
- ◆ **Agent teams as products**: developers publish a team (roles, mandates, plan) and earn a team fee on every released mission.

### MARKETPLACE FLYWHEEL



### SUPPLY-SIDE EXAMPLES: WHAT SELLERS COULD LIST (ILLUSTRATIVE, TO RECRUIT)

| SELLER     | LISTING                                                                    | KIND    | HOW THEY GET PAID                      |
|------------|----------------------------------------------------------------------------|---------|----------------------------------------|
| Developer  | Competitor-pricing agent team: researcher + writer with mandates prefilled | TEAM    | 3 USDC team fee per released mission   |
| Architect  | Code-compliance review of a floor plan, delivered as a marked-up PDF       | SERVICE | 25 USDC from escrow on delivery hash   |
| Researcher | Sector report: EU battery supply chain, 2026 edition                       | DATA    | 12 USDC, graded A to D by the assessor |
| Data owner | Weekly SaaS pricing feed for 500 vendors                                   | API     | 0.05 USDC per call over x402           |

# The model we are testing: *a small fee on every settled mission and agent payment.*

SANDBOX

\$0

Devnet, test USDC, MCP server, demo missions. For builders evaluating Fiducia from Claude Code.

BUILDER

1.5%


of settled agent spend, no monthly fee. Mainnet missions with low caps. Sellers pay nothing.

AGENCY

\$149<sub>/mo</sub>

+ 1% of settled spend. Client workspaces, mandate templates, per-client spend export, priority verifier.

Pricing is a hypothesis to test with design partners, not a published price list. Benchmarks: card processing is about 2.9% + \$0.30 per payment; Fiducia's fee must stay below what an agency would accept as a client pass-through.



WHO PAYS

The buyer, on settlement. Agencies pass it through to the client as a line item next to the data cost.

WHY THEY PAY


A mandate and on-chain receipt per client replaces shared keys, card float and month-end reconciliation.

WHY NOT CARDS


Card fees make sub-dollar agent purchases uneconomic; Solana settles them for under a tenth of a cent.

EXPANSION REVENUE

Team fees on published agent teams, verifier and assessor services, enterprise policy packs.



# From devnet demo to paid agencies in five phases.



## DISTRIBUTION CHANNELS, RANKED BY EXPECTED COST PER PILOT

| # | CHANNEL                                                                | WHY IT WORKS FOR US                                                                 |
|---|------------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| 1 | <b>MCP server</b> in Claude Code, Claude Desktop and other MCP clients | Builders hire a team or buy data from where they already work; one install command. |
| 2 | <b>MCP directories</b> and the official registry                       | Discovery by agent builders searching for payment and data tools.                   |
| 3 | <b>Direct outreach</b> to automation and research agencies             | Short sales cycle; the founder is also the developer.                               |
| 4 | <b>Solana ecosystem:</b> Colosseum, Superteam, x402 builders           | Grants, distribution and early sellers already holding USDC.                        |
| 5 | <b>Proof content:</b> every demo links to Explorer                     | Anyone can verify a refusal or payout themselves.                                   |

## PUBLIC LAUNCH WEEK PLAYBOOK

|          |                                                                               |
|----------|-------------------------------------------------------------------------------|
| D-14     | Seed 100 listings; 3 design-partner case studies with on-chain receipts.      |
| D-7      | Waitlist opens; MCP install guide and mandate templates published.            |
| D-0      | Launch post, Show HN, X thread with live refusal demo, Solana community call. |
| D+1 .. 7 | Daily office hours; first 20 agencies onboarded by hand.                      |
| D+30     | Publish metrics: missions, refusals, settled spend, cap breaches (target 0).  |

## FUNNEL TARGETS FOR THE DESIGN-PARTNER PHASE



## DESIGN-PARTNER OFFER

- ◆ 3 months free on devnet, then capped mainnet at cost.
- ◆ We set up their first mandate templates and seller list by hand.
- ◆ Direct line to the founders; fixes shipped weekly.

## WHAT WE ASK IN RETURN

- ◆ One real client workflow run through Fiducia every week.
- ◆ A 30-minute feedback call every two weeks.
- ◆ Permission to publish an anonymised case study with on-chain receipts.

Where we are *today*.

What we *measure*.

| ITEM                                                          | STATUS          |
|---------------------------------------------------------------|-----------------|
| Solana program: mandates, stages, escrow, refusal, settlement | LIVE · DEVNET   |
| Web app with one-click judge demo                             | LIVE            |
| MCP server (17 tools) incl. hire_team , buy , demand_board    | WORKING         |
| Seller pipeline: grade, price, PII scan, sealed delivery      | DEVNET          |
| Real devnet mission with refusal and payout on Explorer       | VERIFIED        |
| Agents in the demo                                            | SIMULATED AI    |
| Live Claude agents in mission slots                           | NEXT            |
| Design partners, paying customers, revenue                    | NONE YET        |
| Audit, mainnet, real USDC                                     | PLANNED Q1 2027 |

| METRIC                                 | PHASE 1<br>TARGET | PHASE 3<br>TARGET |
|----------------------------------------|-------------------|-------------------|
| Active agencies (weekly)               | 5                 | 40 paying         |
| Missions completed                     | 50                | 1,000 / month     |
| Governed spend settled                 | devnet only       | \$160k / month    |
| Cap breaches                           | 0                 | 0                 |
| Refusals explained to the user         | 100%              | 100%              |
| Listings with a grade                  | 30                | 100+              |
| Disputes resolved inside review window | n/a               | > 95%             |
| Time from install to first mission     | < 15 min          | < 5 min           |

Targets only.

RISKS AND HOW WE REDUCE THEM

| RISK                                                                 | LIKELIHOOD    | MITIGATION                                                                                                             |
|----------------------------------------------------------------------|---------------|------------------------------------------------------------------------------------------------------------------------|
| Agencies will not touch crypto wallets                               | high          | One-click demo buyer today; embedded wallet and card-funded USDC on the roadmap; users see dollars, not tokens.        |
| Thin seller supply, empty catalogue                                  | high          | Seed 30 then 100 listings by hand; demand board; zero seller fees; agent teams as products.                            |
| Big platforms ship spending controls (Stripe, Google, card networks) | medium        | Compose with their rails; our edge is chain-enforced mandates, delivery escrow and per-client receipts across sellers. |
| Smart contract bug                                                   | medium        | Audit before mainnet, multisig upgrades, low per-mission ceilings, binary verification published.                      |
| Agent quality, not spending, is the real blocker                     | medium        | Stage approvals on plan hashes and paid-on-delivery escrow make bad output cheap to reject.                            |
| Regulation of stablecoin payments                                    | low to medium | Non-custodial program vaults; USDC from a regulated issuer; legal review before mainnet in each launch market.         |

LANDSCAPE: RAILS VERSUS CONTROL LAYER

PAYMENT RAILS

x402, card networks' agent programs, Stripe/OpenAI ACP, Google AP2. **They move money.** Fiducia sits on top and decides if a payment may happen.

AGENT WALLETS

Embedded wallets with server-side policies. **Limits enforced by a company's server**, not by the chain.

HACKATHON PEERS

Cordon, OpenBook, OnchainRouter, Turnstile (EVM and Hedera). **None on Solana**; none combine mandates, stages and delivery escrow.

FIDUCIA

Chain-enforced mandate per agent, human approval of the exact plan, escrow paid on delivery, per-client receipts.

# What comes *after the wedge*.

2026 · NOW

RESEARCH AGENTS

Agencies buying data, APIs, reports and expert work per client.

2027

CODING AGENTS

Paying for APIs, compute and code review; delivery = diff and passing build hash.

2027 TO 2028

BOOKING AGENTS

Travel and service bookings inside company policy, stage approval above a threshold.

2028+

AGENT COMMERCE

Agent-to-agent markets; McKinsey sizes agentic commerce at \$3T to \$5T by 2030.

FACTS WE CAN SHOW TODAY

- ◆ Program CfD43mq2...NGSV live on Solana devnet; deployed binary verified against the committed build.
- ◆ Real devnet mission: 2.000001 USDC refused by the chain (OverPerTxCap), 1.00 USDC paid to the approved seller, product hash delivered, fee released.
- ◆ MCP server with 17 tools; Claude Code calls `hire_team` and gets a human approval link.
- ◆ Public repo with 403 test cases across 5 packages.

HYPOTHESES WE STILL HAVE TO PROVE

- ◆ Research agencies feel the accountability pain strongly enough to switch.
- ◆ A 1 to 1.5% settlement fee plus \$149 per month is acceptable as a client pass-through.
- ◆ Sellers will list for agent buyers without a commission incentive.
- ◆ Live model agents keep plans inside mandates often enough to be useful.

SOURCES

|                                                                                                                                                         |                                                                                                                           |
|---------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| Gartner, "Gartner Predicts 2025: AI agents" and press release on agentic AI in enterprise software (Oct 2024)                                           | <a href="https://gartner.com/en/newsroom">gartner.com/en/newsroom</a>                                                     |
| Gartner, "Over 40% of agentic AI projects will be canceled by end of 2027" (25 Jun 2025)                                                                | <a href="https://gartner.com/en/newsroom/press-releases/2025-06-25">gartner.com/en/newsroom/press-releases/2025-06-25</a> |
| Stack Overflow Developer Survey 2025, AI section                                                                                                        | <a href="https://survey.stackoverflow.co/2025/ai">survey.stackoverflow.co/2025/ai</a>                                     |
| McKinsey, "The agentic commerce opportunity" (Oct 2025)                                                                                                 | <a href="https://mckinsey.com">mckinsey.com</a>                                                                           |
| ESOMAR, Global Market Research 2024 report                                                                                                              | <a href="https://esomar.org">esomar.org</a>                                                                               |
| U.S. Census Bureau, 2022 Economic Census First Look (2024)                                                                                              | <a href="https://census.gov/newsroom/press-releases/2024">census.gov/newsroom/press-releases/2024</a>                     |
| Upwork Inc., Form 10-K for FY2025 (gross services volume)                                                                                               | <a href="https://sec.gov">sec.gov</a> · CIK 1627475                                                                       |
| Blockworks Research, State of Solana Q4 2025 (fees, stablecoin supply)                                                                                  | <a href="https://blockworks.co/research">blockworks.co/research</a>                                                       |
| Stripe, US standard pricing (2.9% + 30¢ per card charge)                                                                                                | <a href="https://stripe.com/pricing">stripe.com/pricing</a>                                                               |
| Coinbase x402 launch (May 2025); Google AP2 (Sep 2025); OpenAI and Stripe ACP (Sep 2025); Visa Intelligent Commerce and Mastercard Agent Pay (Apr 2025) | <a href="#">company announcements</a>                                                                                     |
| Hackathon analysis: ETHGlobal ETHOnline 2026 showcase and finale; Superteam × WHU 2026 submissions                                                      | <a href="#">Analysis/ in repo</a>                                                                                         |

Fiducia controls how *AI agents spend money*.

fiducia-orpin.vercel.app

[github.com/MuLaydm10/solana-hackathon](https://github.com/MuLaydm10/solana-hackathon)

Built by Dhruv and Vedant